

Relational Algebra

IDATG2204 Data Modelling and Database Systems

Where are We Now?

- W02: Introduction, Relational Algebra
- W03: SQL
- W04: SQL, Conceptual Modelling
- W05: Conceptual Modelling
- W06: Normalisation
- W07: Logical Modelling, NOSQL
- W08: DB Application Development
- W09: DB Security, Project Kick-off
- W10-W14: Project Work with Peer Review
- W15: Indexing, query processing, concurrency
- W16: Recovery
- W17: More SQL and NOSQL
- W18: Review and Wrap-up

Why Relation Algebra

- Help us understand typical relational query operations without worrying about the SQL syntax
- Relational algebra is important in the implementation of RDBMS'es

Sample Relations

car

<i>id</i>	<i>make</i>	<i>model</i>	<i>year</i>	<i>mileage</i>	<i>fuel</i>	<i>type</i>	<i>price</i>	<i>dealer_id</i>
1	Volkswagen	Passat	2017	97805	diesel	station wagon	425000	Bdf
2	Mazda	CX-3	2019	19777	petrol	suv	378900	Gjvk
3	Volkswagen	UP!	2017	16551	electric	hatchback	125000	Bdf
4	Toyota	RAV4	2019	39661	hybrid	suv	428900	Hmr
5	Mercedes-Benz	C Class	2004	301204	diesel	sedan	31707	Hmr
6	Audi	Q3	2020	18516	diesel	suv	624900	Hrst

dealer

<i>id</i>	<i>city</i>	<i>county_no</i>
Bdf	Bardufoss	54
Bo	Bodø	18
Elv	Elverum	34

county

<i>no</i>	<i>name</i>
30	Viken
3	Oslo
34	Innlandet

Outline

- Relational algebra operations:
 - Unary operations
 - Set and join operations
 - Aggregation and grouping operations
 - (Division operation)

Unary Operations

- A unary operation defines a new relation based on one input relation
 - Projection
 - Selection
 - Rename operation

Projection

- Defines a relation that contains a subset of columns from the input relation
- Notation:
 - $\Pi_{a_1, \dots, a_n}(R)$
- Examples:
 - Car id, make, model, year only:
 - $\Pi_{id, make, model, year}(car)$
 - Car id, year, mileage, price only:
 - $\Pi_{id, year, mileage, price}(car)$
 - Car id, make, dealer_id only:
 - $\Pi_{id, make, dealer_id}(car)$

Selection

- Defines a relation that contains a subset of rows from the input relation that satisfy the specified condition
- Notation:
 - $\sigma_{\text{predicate}}(R)$
- Examples:
 - Mazda cars only:
 - $\sigma_{\text{make}='Mazda'}(\text{car})$
 - No cars older than 2018:
 - $\sigma_{\text{year} \geq 2018}(\text{car})$
 - No cars older than 2018 and mileage larger than 20,000 km:
 - $\sigma_{\text{year} \geq 2018 \text{ AND } \text{mileage} \leq 20000}(\text{car})$

Select conditions

- Basic conditions:
 - Comparison: =, <>, <, <=, >, >=
 - Range: BETWEEN, NOT BETWEEN
 - `year BETWEEN 2013 AND 2015`
 - Set membership: IN, NOT IN
 - `county_name IN ('Viken', 'Oslo', 'Innlandet')`
 - Pattern match: LIKE, NOT LIKE
 - `model LIKE 'CX%'`
 - Null: IS NULL, IS NOT NULL
 - `price IS NULL`
- Logical expression:
 - NOT, AND, OR
 - Combined with parenthesis, if needed

Combining Selections and Projections

- Examples:
 - Car id, make, model, and year for cars older than 2010:
 - $\Pi_{id, make, model, year}(\sigma_{year < 2010}(car))$
 - which can be decomposed to:
 - $R \leftarrow \sigma_{year < 2010}(car)$
 $\Pi_{id, make, model, year}(R)$
 - which is equivalent to:
 - $\sigma_{year < 2010}(\Pi_{id, make, model, year}(car))$
 - But:
 - $\Pi_{id, make, model}(\sigma_{year < 2010}(car))$
 - is not equivalent to:
 - $\sigma_{year < 2010}(\Pi_{id, make, model}(car))$

(Rename Operation)

- Provides a new name for a relation - and optionally the attribute names:
- Notation:
 - $\rho_{S(a_1, \dots, a_n)}(R)$
- Example:
 - Rename the county relation to Norwegian
 - $\rho_{fylke(nr, navn)}(county)$

Outline

- Relational algebra operations:
 - Unary operations
 - Set and join operations
 - Aggregation and grouping operations
 - (Division operation)

Set and Join Operations

- A set and join operation defines a new relation based on two input relations
 - Union
 - (Difference)
 - (Intersection)
 - Cartesian product
 - Equijoin
 - (Theta join)
 - (Natural join)
 - (Semijoin)
 - Outer join

Union

- Defines a relation containing all the tuples of the two input relations; the two input relations must be union-compatible:
 - Same number of attributes
 - Corresponding attributes from the same domain
- Notation:
 - $R \cup S$
- Example:
 - Names of all counties, dealer cities, and car makers
 - $\Pi_{\text{name}}(\text{county}) \cup \Pi_{\text{city}}(\text{dealer}) \cup \Pi_{\text{make}}(\text{car})$

Set and Join Operations

- A set and join operation defines a new relation based on two input relations
 - Union
 - (Difference)
 - (Intersection)
 - Cartesian product
 - Equijoin
 - (Theta join)
 - (Natural join)
 - (Semijoin)
 - Outer join

Cartesian Product

- Defines a relation that is the concatenation of every single tuple in one input relation with every single tuple of the other
- Notation:
 - $R \times S$
- Example:
 - All combinations of county name and city for county number 18 and 30:
 - $\Pi_{\text{name}}(\sigma_{\text{no IN (18,30)}}(\text{county})) \times \Pi_{\text{city}}(\sigma_{\text{county_no IN (18,30)}}(\text{dealer}))$
 - Decomposed:
 - $R \leftarrow \Pi_{\text{name}}(\sigma_{\text{no IN (18,30)}}(\text{county}))$
 $S \leftarrow \Pi_{\text{city}}(\sigma_{\text{county_no IN (18,30)}}(\text{dealer}))$
 $R \times S$

Equijoin

- Defines a relation that contains tuples satisfying the equality predicate F from the Cartesian product of the two input relations
- Notation:
 - $R \bowtie_F S$
- Example:
 - Car and corresponding dealer information:
 - $\text{car} \bowtie_{\text{dealer_id} = \text{dealer.id}} \text{dealer}$
 - Dealer and corresponding county information:
 - $\text{dealer} \bowtie_{\text{county_no} = \text{no}} \text{county}$
 - Car, dealer, and county information joined:
 - $\text{car} \bowtie_{\text{dealer_id} = \text{dealer.id}} \text{dealer} \bowtie_{\text{county_no} = \text{no}} \text{county}$

Set and Join Operations

- A set and join operation defines a new relation based on two input relations
 - Union
 - (Difference)
 - (Intersection)
 - Cartesian product
 - Equijoin
 - (Theta join)
 - (Natural join)
 - (Semijoin)
 - Outer join

Outer join

- Defines a join where tuples in one input relation that do not match tuples in the other are also included
- Notation:
 - $R \bowtie_F S$ (left outer join)
 - $R \ltimes_F S$ (right outer join)
- Example:
 - All counties including dealer cities where these exist:
 - $\text{county} \bowtie_{\text{county_no} = \text{no}} \text{dealer}$
 - All cars and dealers, including dealers that have no cars:
 - $\text{car} \ltimes_{\text{dealer_id} = \text{dealer.id}} \text{dealer}$
 - All counties, cities, and corresponding cars:
 - $\text{county} \bowtie_{\text{county_no} = \text{no}} \text{dealer} \bowtie_{\text{dealer_id} = \text{dealer.id}} \text{car}$

Outline

- Relational algebra operations:
 - Unary operations
 - Set and join operations
 - Aggregation and grouping operations
 - (Division operation)

Aggregation

- Defines a relation that contains one tuple holding the value(s) of the listed aggregate function(s)
- Notation:
 - $_AL(R)$
- Example:
 - Number of dealers in the company:
 - `COUNT id(dealer)`
 - Result column renamed to `dealerCount`:
 - `$\rho_{R(dealerCount)}(COUNT\ id(dealer))$`
 - Min, max, and average price of a Volkswagen car:
 - `MIN price, MAX price, AVG price($\sigma_{make='Volkswagen'}(car)$)`

Aggregation Functions

- Common ones are:
 - COUNT -- Count number of values
 - SUM -- Sum of the values
 - AVG -- Average of the values
 - MIN -- Smallest value
 - MAX -- Largest value

Grouping

- Groups the tuples of a relation by one or more attributes and applies aggregate function(s) to each group.
- Notation:
 - $GA_{AL}(R)$
- Example:
 - Number of cars of each make:
 - `make COUNT id(car)`
 - Min, max, and average price per car make and model:
 - `make, model MIN price, MAX price, AVG price(car)`
 - Number of dealers per county - for ALL counties:
 - `name COUNT id (county ⋈county_no = no dealer)`