

SQL

IDATG2204 Data Modelling and Database Systems

Where are We Now?

- W02: Introduction, Relational Algebra
- **W03: SQL**
- W04: SQL, Conceptual Modelling
- W05: Conceptual Modelling
- W06: Normalisation
- W07: Logical Modelling, NOSQL
- W08: DB Application Development
- W09: DB Security, Project Kick-off
- W10-W14: Project Work with Peer Review
- W15: Indexing, query processing, concurrency
- W16: Recovery
- W17: More SQL and NOSQL
- W18: Review and Wrap-up

Outline

- Intro to SQL
- SQL SELECT:
 - SELECT, DISTINCT, ORDER BY
 - FROM
 - WHERE
 - Subqueries
 - GROUP BY
 - HAVING
 - Union

Intro to SQL

- SQL is a **declarative** database query language:
 - Describing what the result would look like
 - Leave to the DBMS to decide how to create the result
- SQL reserved words are typically written in upper case
- Database languages are divided in two sub-languages:
 - Data Manipulation Language (DML)
 - For modifying user data: Create, Retrieve, Update, Delete (CRUD)
 - `SELECT`, `INSERT`, `UPDATE`, `DELETE`
 - Data Definition Language (DDL)
 - For modifying database catalogue/data dictionary
 - `CREATE`, `ALTER`, `DROP`

Intro to SQL SELECT

- General form:

```
SELECT      [DISTINCT] { * | [columnExpression [alias]]  
                                     [, ...] }  
  
FROM        TableName [alias] [[, ...] | [joinExpression]]  
[WHERE      condition]  
[GROUP BY  columnList] [HAVING condition]  
[ORDER BY  columnList]
```

- Basic sequence of processing:

1. FROM
2. WHERE
3. GROUP BY
4. HAVING
5. SELECT
6. ORDER BY

Outline

- Intro to SQL
- **SQL SELECT:**
 - SELECT, DISTINCT, ORDER BY
 - FROM
 - WHERE
 - Subqueries
 - GROUP BY
 - HAVING
 - Union

SQL SELECT (1)

- General form:
 - `SELECT [DISTINCT] { * | [columnExpression [alias]]
[, ...] }`
 - What relational algebra operation is this equal to?
 - The projection operation (π): $\Pi_{a_1, \dots, a_n}(R)$
- Wildcard (`*`):
 - For retrieving all columns in the relation
 - Example: Find all cars for sale:
 - `SELECT * FROM cars`
- `columnExpression` may be a list of columns:
 - Example: Find the makes and the models of cars for sale:
 - `SELECT make, model FROM car`

SQL SELECT (2)

- General form:

- `SELECT [DISTINCT] { * | [columnExpression [alias]]
[, ...] }`

- `columnExpression` may also be:

- Calculated fields, e.g.:

- `SELECT CONCAT(UPPER(make), ' ', UPPER(model))`
 - `SELECT MIN(mileage)`

- `alias`:

- For renaming some of the attributes in the result relation

- `SELECT MIN(mileage) AS smallest FROM car`

- What relational algebra operation is this equal to?

- The rename operation: $\rho_{S(a_1, \dots, a_n)}(R)$

SQL SELECT (3)

- General form:

- `SELECT [DISTINCT] { * | [columnExpression [alias]]
[, ...] }`

- DISTINCT:

- When we want duplicates removed, e.g.,:

- `SELECT DISTINCT year FROM car`

- Example:

- List unique car and model names, renaming columns to Norwegian:

- `SELECT DISTINCT make AS merke, model AS modell`

ORDER BY

- Data is to be considered unordered in an RDBMS:
 - The RDBMS decides on the order
- We may ask the RDBMS to return tuples in a specific order:
 - `ORDER BY column [DESC] [, ...]`
- Example
 - Order car tuples on make, model, and year:
 - `ORDER BY make, model, model_year`
 - Order car tuples on year (descending), make, model:
 - `ORDER BY model_year DESC, make, model`

Outline

- Intro to SQL
- SQL SELECT:
 - SELECT, DISTINCT, ORDER BY
 - **FROM**
 - WHERE
 - Subqueries
 - GROUP BY
 - HAVING
 - Union

SQL FROM

- The FROM clause is where we define what table to query:
 - `SELECT * FROM car`
- or what tables to combine data from and how to join:
 - INNER JOIN - for equijoins in the relational algebra:
 $R \bowtie_F S$
 - LEFT/RIGHT OUTER JOIN - as for the relational algebra:
 $R \bowtie_F^L S$
 $R \bowtie_F^R S$

SQL INNER JOIN

- Example:
 - Find car information together with dealer city and county name:
 - ```
SELECT car.*, city, name FROM car
INNER JOIN dealer ON dealer_id = dealer.id
INNER JOIN county ON county_no = no
```

# SQL OUTER JOIN

- Example:
  - Find car information along dealer city and county name, for all cities and counties even when there are no cars:
    - ```
SELECT car.*, city, name FROM car
  RIGHT OUTER JOIN dealer ON dealer_id = dealer.id
  RIGHT OUTER JOIN county ON county_no = no
```

Outline

- Intro to SQL
- SQL SELECT:
 - SELECT, DISTINCT, ORDER BY
 - FROM
 - **WHERE**
 - Subqueries
 - GROUP BY
 - HAVING
 - Union

SQL WHERE (1)

- What relational algebra operation is this equal to?
 - The selection operation: $\sigma_F(R)$
 - Five basic search conditions/predicates (repeated):
 - Comparison (`=`, `<>`, `<`, ...)
 - Range (`[NOT] BETWEEN`)
 - Set membership (`[NOT] IN (...)`)
 - Pattern match (`[NOT] LIKE '...'`)
 - Null (`IS [NOT] NULL`)
 - Boolean expressions (repeated):
 - `AND`, `OR`, `NOT`, `(,)`

SQL WHERE (2)

- MariaDB has many built-in functions:
 - <https://mariadb.com/kb/en/built-in-functions/>
- Example: Find cars from 2017-2019 where mileage is no more than 40000 and fuel type is not electric or hybrid:
 - ```
SELECT * FROM car
WHERE model_year BETWEEN 2017 AND 2019
 AND mileage <= 40000
 AND fuel NOT IN ('electric', 'hybrid')
```

# Outline

- Intro to SQL
- SQL SELECT:
  - SELECT, DISTINCT, ORDER BY
  - FROM
  - WHERE
  - **Subqueries**
  - GROUP BY
    - HAVING
  - Union

# Introduction to Subqueries

- Types of subqueries:
  - Scalar subqueries:
    - Returning scalar values,  
e.g.: `SELECT AVG(price) FROM car`
  - (Row subquery)
  - Table subqueries:
    - Returning a table,  
e.g., `SELECT DISTINCT county_no FROM dealer`
  - Boolean subqueries:
    - `EXISTS/NOT EXISTS`

# Aggregation Subquery

- Example:
  - Find all cars of the oldest model:
    - ```
SELECT * FROM car
WHERE model_year =
      (SELECT MIN(model_year) FROM car)
```
 - List the make, model, and year of all cars that are more expensive than the average car price, and show by how much the price is greater than the average:
 - ```
SELECT make, model, model_year, price -
 (SELECT AVG(price) FROM car)
FROM car
WHERE price > (SELECT AVG(price) FROM car)
```

# ALL Subquery

- The ALL keyword may be used with subqueries that produce a single column of numbers
  - The condition is only true if it is satisfied by all values produced
- Example:
  - List the model years for which there are fewest cars for sale:
    - ```
SELECT model_year, COUNT(id) as count
FROM car
GROUP BY model_year
HAVING count <= ALL
(SELECT COUNT(id)
FROM car
GROUP BY model_year)
```

IN/NOT IN Subquery

- We have already used IN/NOT IN:
 - `WHERE fuel NOT IN ('hybrid', 'electric')`
- A similar set may be produced by a subquery
- IN subqueries should usually be replaced by JOIN
- Example:
 - Find the names of counties where there are no dealers:
 - ```
SELECT name
FROM county
WHERE no NOT IN
 (SELECT DISTINCT county_no
 FROM dealer)
```

# (Row Subquery Example)

- Say we are interested in any Volkswagen Passat, Toyota Corolla, or Audi A3:
  - ```
SELECT * FROM car
WHERE (make, model) IN
      (('Volkswagen', 'Passat'),
       ('Toyota', 'Corolla'),
       ('Audi', 'A3'))
```
- If the user preferences are stored in table `user_pref`:
 - ```
SELECT * FROM car
WHERE (make, model) IN
 (SELECT make, model
 FROM user_pref
 WHERE user_id = 'rune.Hjelsvold@ntnu.no')
```

# EXISTS/NOT EXISTS Subquery

- Produce true or false depending on whether the subquery is non-empty or empty, respectively
- The subquery "sees" attributes in the outer query
- Is usually more efficient than equivalent `IN/NOT IN` subquery
- Example:
  - Find the names of counties where there are no dealers:
    - ```
SELECT name
FROM county
WHERE NOT EXISTS
      (SELECT *
       FROM dealer
       WHERE no = county_no)
```


Outline

- Intro to SQL
- SQL SELECT:
 - SELECT, DISTINCT, ORDER BY
 - FROM
 - WHERE
 - Subqueries
 - **GROUP BY**
 - **HAVING**
 - Union

SQL GROUP BY (1)

- What relational algebra operation is this equal to?
 - The grouping operation: $\text{GA}_{\text{AL}}(R)$
- The attributes returned in a group by query may only be:
 - Grouping attributes
 - Attributes with only one value per group
 - Aggregate values
- Example:
 - Find the min, max, and average price of cars per dealer (city):
 - ```
SELECT city, MIN price, MAX price, AVG price
FROM car
 INNER JOIN dealer ON dealer_id = dealer.id
GROUP BY dealer.id
```

# SQL GROUP BY (2)

- Aggregation variations:
  - Aggregating distinct values only:
    - `COUNT (DISTINCT model_year)`
  - Counting number of defined values:
    - `COUNT (comment)`
  - Counting number of tuples:
    - `COUNT (*)`

# SQL HAVING

- The WHERE clause identifies tuples to be grouped:
  - $GA_{AL}(\sigma_F(R))$
- The HAVING clause identifies what groups to return:
  - $\sigma_{F_{AL}}(GA_{AL}(R))$
- Example:
  - Find the number of cars of each make and model, but only when there are more than 2 cars of that make and model:
    - ```
SELECT make, model, COUNT(id) AS car_count
FROM car
GROUP BY make, model
HAVING car_count > 2
```

Outline

- Intro to SQL
- SQL SELECT:
 - SELECT, DISTINCT, ORDER BY
 - FROM
 - WHERE
 - Subqueries
 - GROUP BY
 - HAVING
 - Union

SQL UNION

- The UNION can be used to merge different query results:
 - $R \cup S$
- Example:
 - Find the names of counties, dealer cities, and car makers:
 - ```
SELECT name FROM county
UNION
SELECT city FROM dealer
UNION
SELECT make FROM car
```